

Dr. Mark AI Clone: Strategy

Mrwan Awad, AI Digital Clone Developer Assignment, September 2026

Capturing how he thinks and talks

Four inputs. **Knowledge:** course outlines, slides, posts and transcripts of his existing talks, via Whisper. **Style:** a short written style guide plus 20 to 30 real samples of his writing. **Decision-making:** one 90-minute interview built on decision questions, not biography, such as "when would you turn a student away" and "what claims do you never make". His answers become written rules. **Corrections:** every edit he makes to a draft is saved and injected into every later prompt. The clone improves because he edits it, not because it is retrained, which keeps his time cost low and the behaviour auditable.

When the clone does not know

Unconfirmed facts live in one file, where null means *not decided*. The model is told never to guess, estimate, or give a typical range for a null value. When a fact is missing it says so plainly, offers to follow up, and writes the question to a queue for Dr. Mark. Price, dates, venue, clinical claims, patient cases and anything unapproved always escalate, even when an answer looks available. This is the part most clone builds get wrong: a clone that invents a price is worse than no clone, because the error arrives in his voice.

Voice and video

Need	Tool	Required from him
Voice	ElevenLabs Professional Voice Clone, \$22/mo	30 min clean audio minimum, 2 to 3 hours for production quality
Video	HeyGen custom avatar, \$29 to \$49/mo	2 to 5 min to camera, plus a mandatory verbal consent clip . Training 2 to 4 hours
Transcripts	Whisper	Existing recordings, no new effort

ElevenLabs Instant Voice Clone needs only 1 to 3 minutes from a phone, so we prototype with it in week 1 and replace it after the studio session.

Configure versus build

Configured: the language model, ElevenLabs, HeyGen, Slack or email for approvals, GitHub and hosting. Built by me: the knowledge base and profile, the never-guess rule and escalation list, the pending questions queue, the corrections loop, the idea-to-video pipeline, and the automated checks. The vendors sell a voice and a face; neither knows Dr. Mark, and neither will refuse to invent a price.

Idea to finished video

Idea → script draft → he approves or edits → voice → video → final check → publish. His edits are saved as corrections and feed the next draft. One human gate, always before publication, never after.

Challenges, limits, dependencies

Output quality tracks input volume; thin material produces generic writing. Video avatars read as slightly off during large hand movements, so talking-head formats work and demonstrations do not. No clinical claim ships without his approval, and the clone is not a clinical advice channel. Vendor dependency is real, so the provider layer is deliberately swappable. No patient data enters any tool. The Colombian audience is Spanish-speaking; English-only is a scope decision.

Timeline

Weeks 1 to 2, interview, gather material, text clone. Week 3, voice. Weeks 4 to 5, video avatar and approval pipeline. Weeks 6 to 7, real use and tuning. Week 8, handoff and account transfer.

Fully functional means: he approves most drafts with light edits; the clone never states a price, date or clinical claim he has not confirmed; voice and video pass his own review; and a video goes from idea to published in under two days with one approval step.

Costs

	One-time	Monthly
Development and setup	\$6,000 to \$7,000	

	One-time	Monthly
ElevenLabs Creator		\$22
HeyGen Creator or Pro		\$29 to \$49
LLM API usage		\$20 to \$100
Hosting and automation		\$20 to \$50
Maintenance, 5 to 10 hours		\$175 to \$350

About \$6,000 to \$7,000 to build, \$250 to \$550 per month to run. HeyGen bills in credits and models bill per token, so the monthly figure tracks how much he publishes: four videos a month and forty are not the same bill. Maintenance is optional after handoff, because he owns the accounts, code and documentation.

What we need from Dr. Mark

Decision interview 1.5 hours, voice recording 1 to 2 hours, video filming including the consent clip 1 to 2 hours, then about 30 minutes a week reviewing drafts through week 7, and about an hour a week on approvals after launch. **Roughly 8 to 10 hours across the first two months**, front-loaded. After that the ask is approvals, which is work he already does.

Privacy, ownership, handoff

Every account (LLM, ElevenLabs, HeyGen, GitHub, hosting) opens **in his name**, billed to him, with me added as a team member, not the reverse. His voice, face and content remain his; written consent before cloning; deletion on request. The code lives in his repository and he owns it outright. No patient data in any tool. At handoff: an operating guide, a full account inventory, and **every key rotated so my access ends**. A clone he cannot run without me is a liability, not an asset.

What I add beyond existing AI tools

Anyone can buy ElevenLabs and HeyGen this afternoon. Neither knows Dr. Mark or has any reason to refuse a question it cannot answer. The work is the layer between them: his knowledge, his decision rules, a clone that recognises the edge of what it knows and routes that to him, and a correction loop that makes it more like him every week, in one workflow he owns rather than four subscriptions he babysits.

Sources reviewed

First Movers: [home](#), [AI Labs](#) (\$250/mo, \$2,500/yr), [blog](#), notably "[The AI Employee](#)". Julia McCoy: "[How I Clone Myself With HeyGen & ElevenLabs](#)", "[Build Your AI Clone Step-by-Step](#)". Vendors: [ElevenLabs pricing](#) and [voice cloning docs](#), [HeyGen pricing](#) and [Digital Twin FAQ](#).

First Movers publicly describes its own builds as HeyGen plus ElevenLabs plus custom Claude training. The stack proposed here matches what they already run, deliberately.

Assumptions

He has existing content to train on; if not, add 1 to 2 weeks for capture. The clone is for marketing and education, not patient care. English is the working language; Spanish for Colombia is roughly one extra week plus a second voice. Development is estimated at about 180 hours at a \$35/hour blended rate. Vendor prices and recording requirements were taken from public pages on 17 September 2026 and can change. First Movers' published figures (300K subscribers, \$3M run-rate) are their own claims, cited as context, not verified by me.